

# DATA SCI 7030: Database and Analytics

Tozammel Hossain  
Oct 29, 2020



Data Science & Analytics  
University of Missouri

# Agenda

- Workflow in a Database Design
  - Requirement analysis (done)
  - The Entity-Relationship Model (done)
  - ERD to DDL

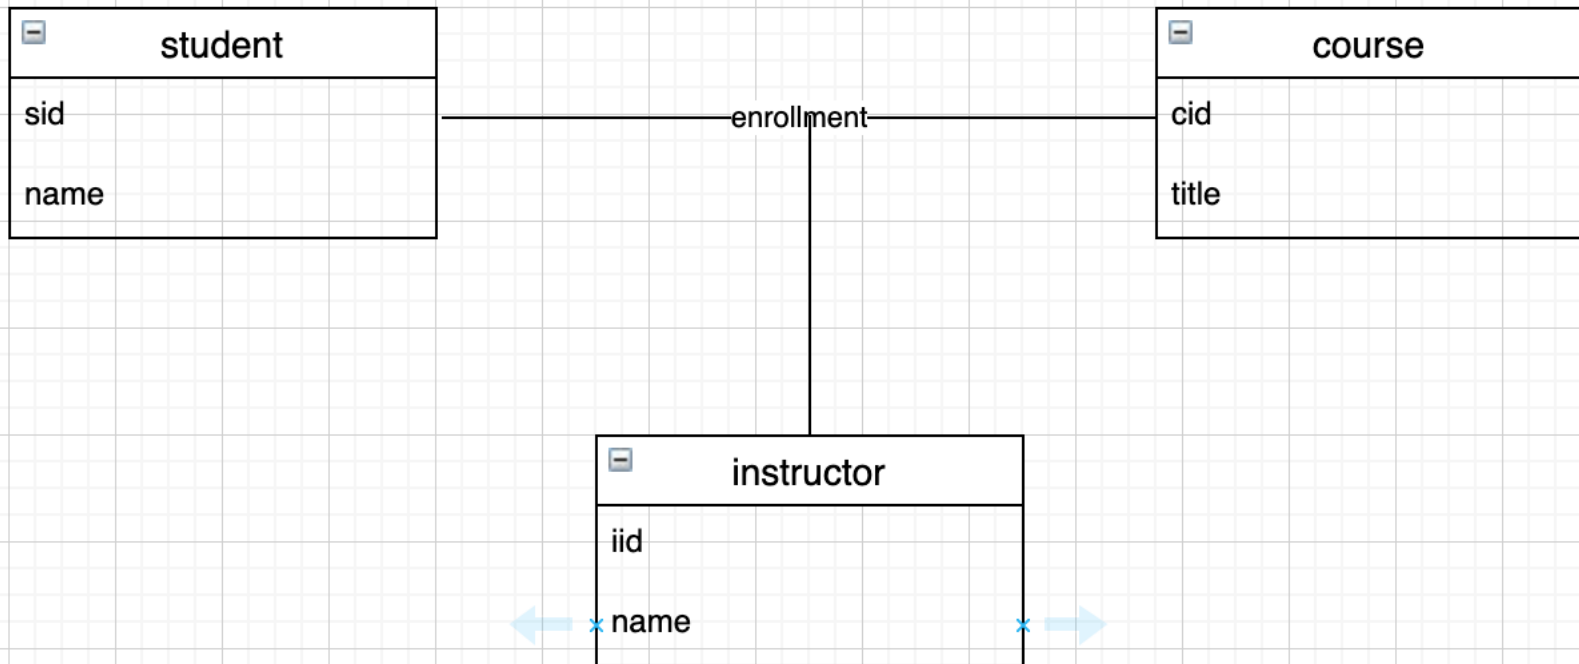
# ER Modeling Example

- Requirements
  - Create a database for a course enrollment system in a university.
  - Students enroll into a courses a semester which are taught by instructors.
  - A course is taught by one and only one instructor

# Entities & Relation

- Student
  - sid, name
- Course
  - cid, title
- Instructor
  - iid, name

# First Conceptual Model



# Limitations

- At the physical/implementation level, RDBMS doesn't support complex relations
  - without introducing overheads/redundancy
- Design principles play role here
  - while updating data, a minimal numbers of table should be involved
  - e.g., previous example
    - keep student/course/instructor info in the same table
      - same info will be repeated again/again
    - Now think about the updating a student?

# Refining ERD

- Entities should participate in a relationship unless we have only one table in the database
- **one-to-one and one-to-many binary relationships** are good
  - we don't need to refine anything here

# Refining ERD

- **many-to-many binary relations** cannot be implemented without introducing redundancy
  - replace the relation with an intermediate entity
    - create the bridge/establish relationship
    - aka **association entity**
  - relate two original entities to this association entity

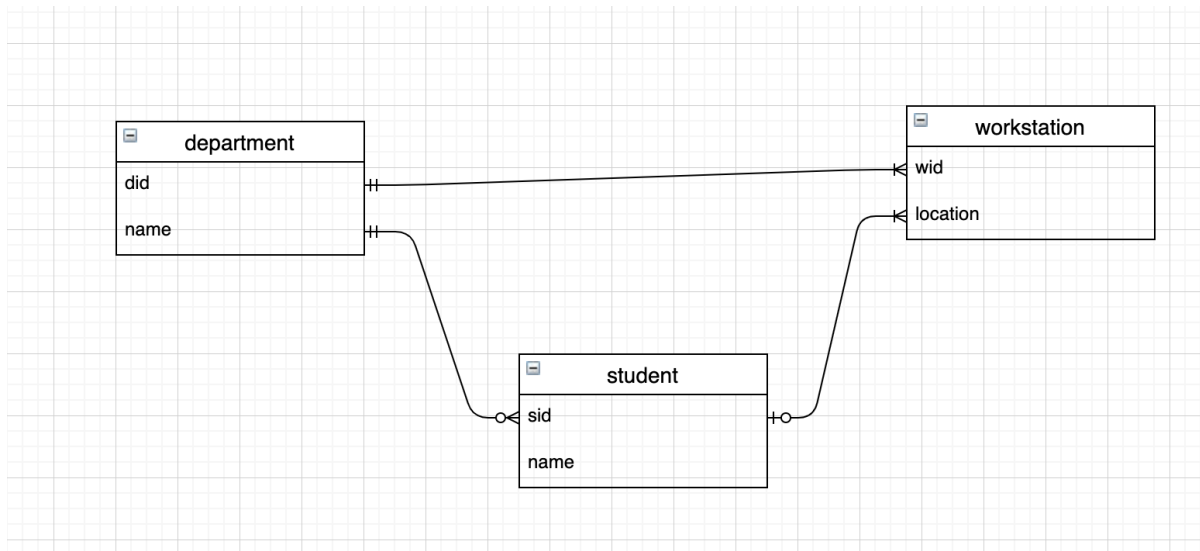


# Refining ERD

- Complex relationship (ternary or n-ary)
  - No way database can implement this without overheads
  - Solution
    - Similar to **many-to-many binary relations**
    - Introduce an **association entity** for the complex relation
    - relate 3/n original entities to this association entity

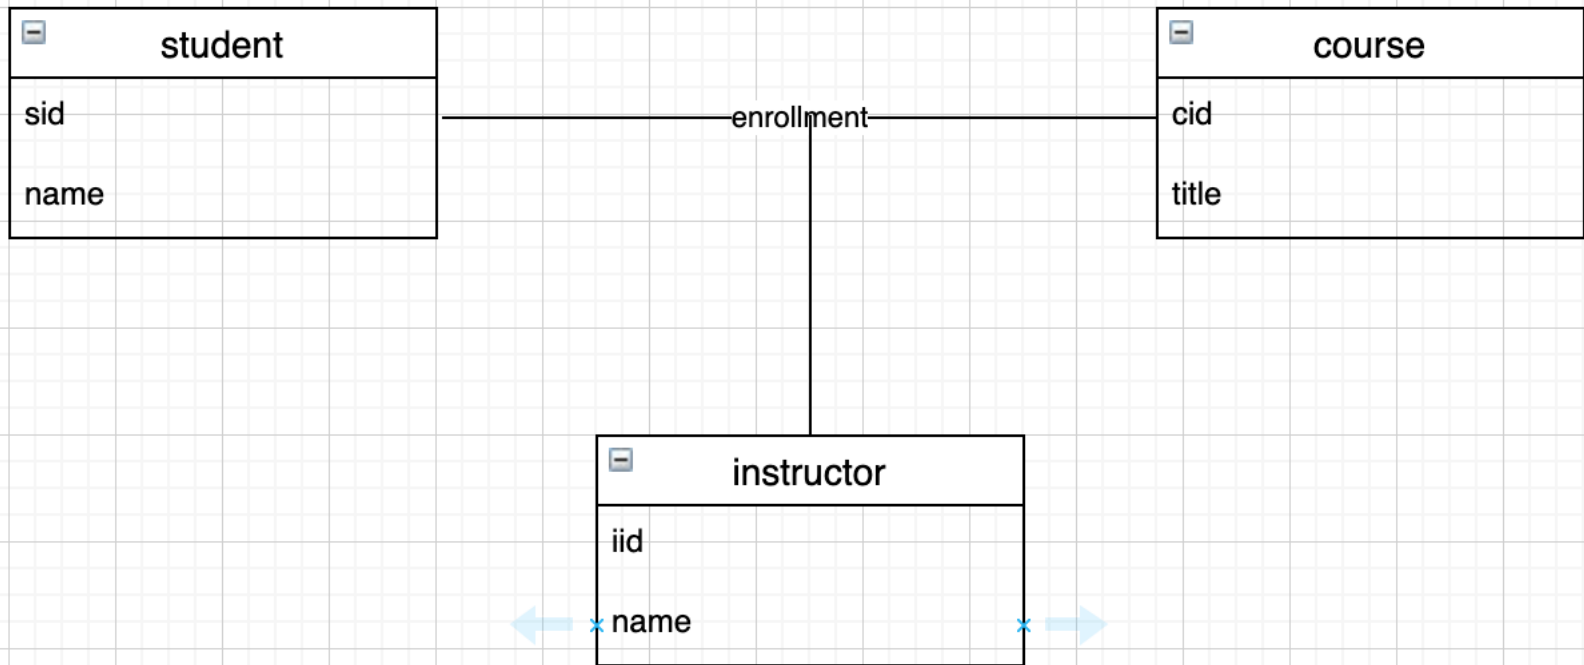
# Refining ERD

- Eliminate redundant relations



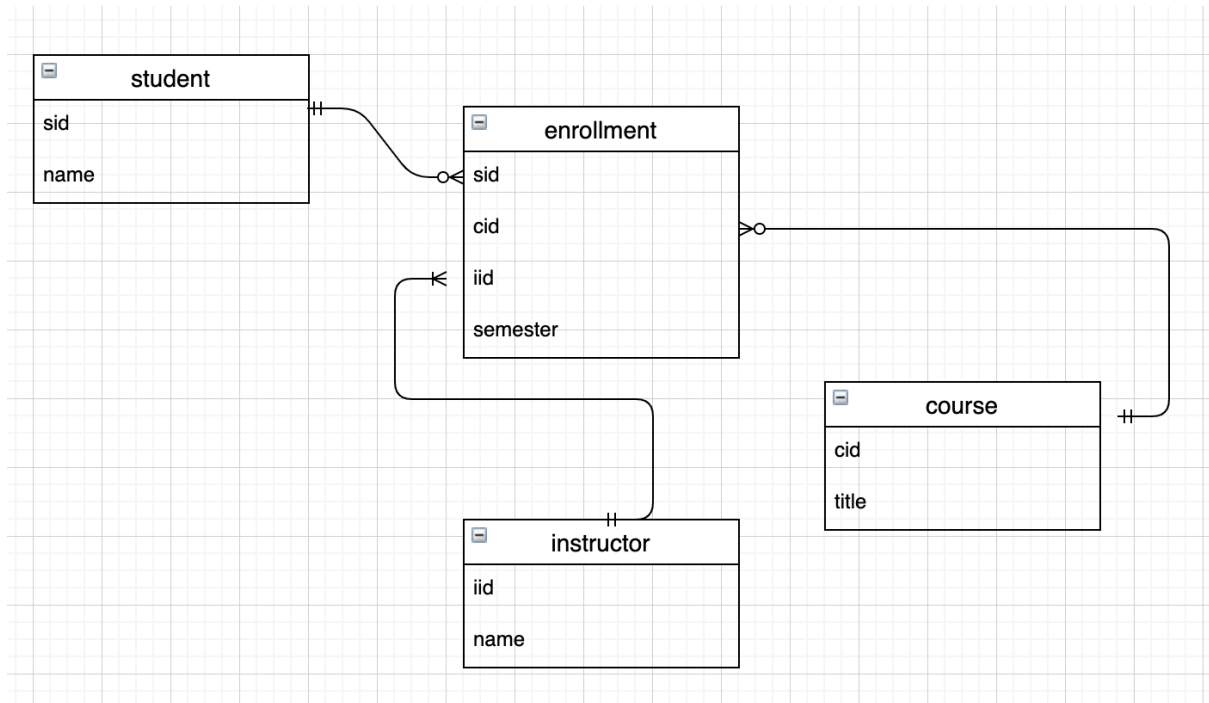
- Relation between department and workstation is redundant
  - Connection between department and workstation could be established via student table

# Refining Course ERD



**Ternary Relationship**

# Final Refinement



- Is the above solution good enough? Is there any other solutions? Yes.
  - Involves database normalization techniques

# ERD to DDL

- Entities and Simple Attributes
  - An entity type within ER diagram is turned into a table
  - The key attribute of the entity is the **primary key** of the table which is usually underlined
    - Uniquely identify a row in the table
  - E.g. simple textual description of schema
    - student(sid,name)
    - cours(cid,title)

# Converting ERD to DDL

- Entities and Simple Attributes
  - Convert to an SQL statement for creating the tables

```
CREATE TABLE STUDENT(  
    sid char(20),  
    name varchar(20)  
);
```

- Note SQL syntax
  - case insensitive (unlike python/r)
  - typed (unlike python/r)
  - semicolon for ending a stmt (unlike python)

# *Using CHAR and VARCHAR*

Value	CHAR(4)	Storage Required	VARCHAR(4)	Storage Required
"	'    '	4 bytes	"	1 byte
'ab'	'ab  '	4 bytes	'ab'	3 bytes
'abcd'	'abcd'	4 bytes	'abcd'	5 bytes
'abcdefgh'	'abcd'	4 bytes	'abcd'	5 bytes

# Activity

- Switch to Jupyter notebook
- Create a dataset for this schema
- No need to fulfill key constraints



# Destroying and Altering Relations

- **DROP TABLE Student;**
  - Destroy the relation student
    - schema is deleted
    - rows (i.e., tuple) are deleted
- **ALTER TABLE Student**  
**ADD DOB date;**
  - The schema of Student is altered by adding a new field
  - every tuple in the current instance is extended with a *null* value in the new field

# Integrity Constraints (ICs)

- Condition that must be true for *any* instance of the database
  - e.g., domain constraint
    - The value of field came from a predefined set
- A **legal** instance of a relation is one that satisfies all specified ICs
  - DBMS should not allow illegal instances
- If the DBMS checks ICs, stored data is more faithful to real-world meaning

# *Where do ICs Come From?*

- ICs are based upon the semantics of the real-world enterprise
  - Discover during requirement analysis
- Key and foreign key ICs are the most common

# *Primary Key Constraints*

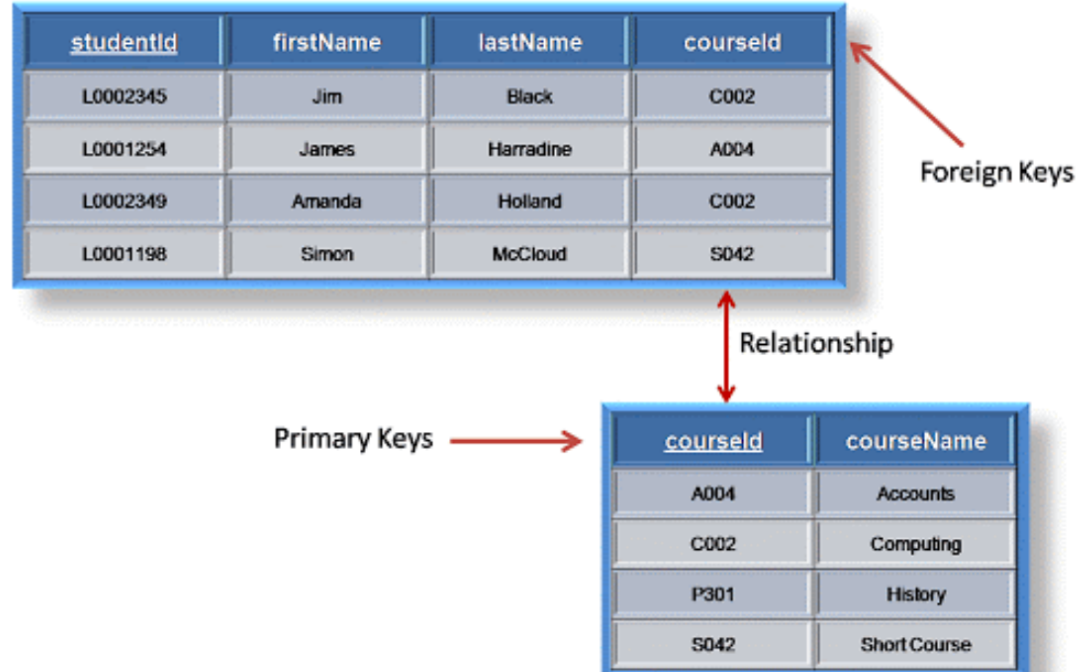
- A set of fields is a *key* for a relation if :
  - No two distinct tuples can have same values in all key fields, and
  - This is not true for any subset of the key.  
(minimal set of attributes)
- E.g., *sid* is a key for Students.
  - What about *name*?
  - What about (sid, name)?
  - The set {*sid*, *gpa*} is a superkey.

# *Primary Key Constraints*

```
CREATE TABLE Student (  
    sid CHAR(20),  
    name VARCHAR(20),  
    PRIMARY KEY(sid)  
);
```

# Foreign Keys, Referential Integrity

- Set of fields in one relation that is used to 'refer' to a tuple in another relation
  - Must correspond to primary key of the second relation



# Foreign Keys, Referential Integrity

- E.g. *sid* is a foreign key in enrollment table referring to Students
- *Foreign Keys in SQL*

```
CREATE TABLE Enrollment (  
    sid CHAR(20),  
    cid CHAR(20),  
    iid CHAR(20),  
    PRIMARY KEY (sid,cid,iid),  
    FOREIGN KEY (sid) REFERENCES Student (sid)  
    ....  
);
```